

Analysis of Algorithms 1 (Fall 2011) Istanbul Technical University Computer Eng. Dept.



Chapter 14 Augmenting Data Structures

Last updated: December 16, 2009

Purpose

Understanding what “augmenting a data structure” means
(augmenting=extending)

Go through examples cases where a known data structure is modified to solve a new problem.

Outline

Augmenting a data structure

Red and Black tree for order statistics
(SELECT and RANK)

Interval trees

Augmenting a Data Structure

It is unusual to have to design an all-new data structure from scratch.

It is more common to take a data structure that you know and store additional information in it.

With the new information, the data structure can support new operations.


But. . . you have to figure out how to ***correctly maintain the new information without loss of efficiency.***

Augment Red-Black Trees

So that we have

- The usual dynamic-set operations
 - **INSERT(S, x)**: inserts element x into set S.
 - **MAXIMUM(S)**: returns element of S with largest key.
 - **EXTRACT-MAX(S)**: removes and returns element of S with largest key.
 - **INCREASE-KEY(S, x, k)**: increases value of element x.s key to k. Assume $k \geq x$.s current key value.
- PLUS the following order statistics related operations:
 - **OS-SELECT(x, i)**: return pointer to node containing the i th smallest key of the subtree rooted at x.
 - **OS-RANK(T, x)**: return the rank of x in the linear order determined by an inorder walk of T .

Red-black trees




BSTs (Binary Search Tree) with an extra one-bit color field in each node.

Red-black properties:

1. Every node is either red or black.
2. The root and leaves (NIL's) are black.
3. If a node is red, then its parent is black.
4. All simple paths from any node x to a descendant leaf have the same number of black nodes = $\text{black-height}(x)$.

Red-black tree example



REMINDER

Red-Black Tree Augmenting

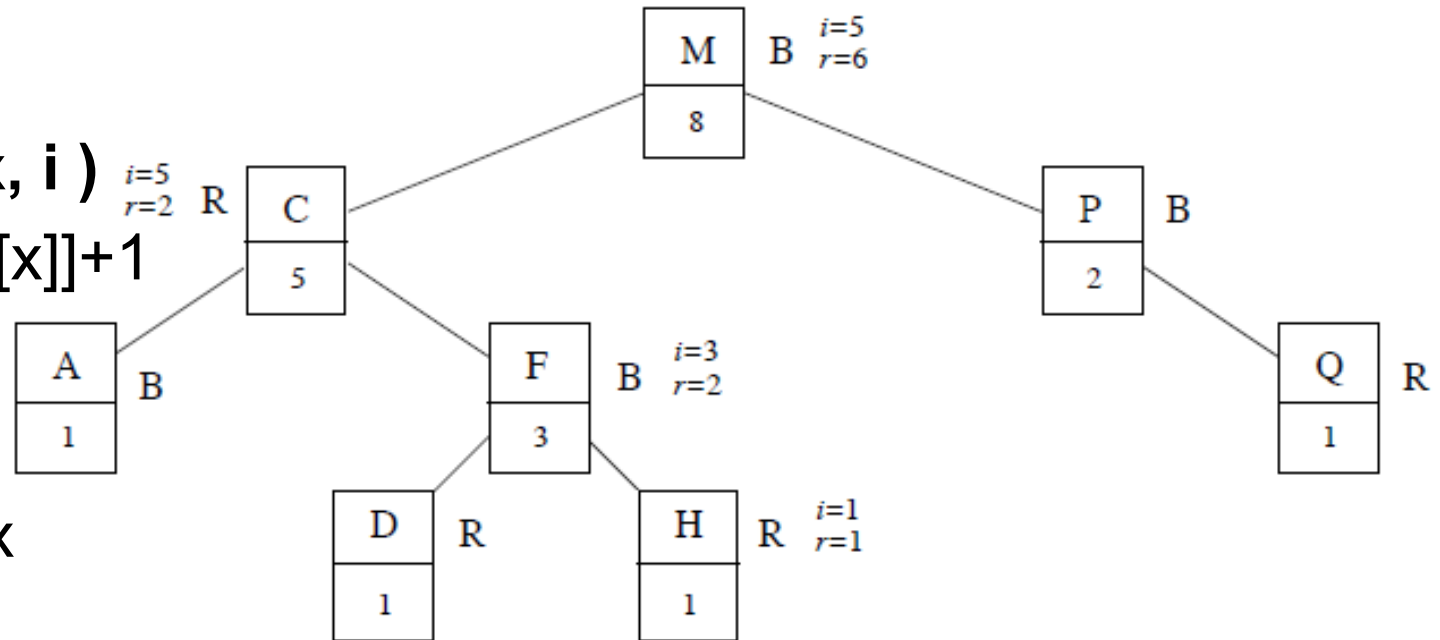
- store in each node x :
- $\text{size}[x] = \#$ of nodes in subtree rooted at x .
 - Includes x itself.
 - Does not include leaves (sentinels).
- Define for sentinel $\text{size}[\text{nil}[T]] = 0$.
- Then $\text{size}[x] = \text{size}[\text{left}[x]] + \text{size}[\text{right}[x]] + 1$
- Note: OK for keys to not be distinct. Rank is defined with respect to position in inorder walk. So if we changed D to C , rank of original C is 2, rank of D changed to C is 3.

OS-SELECT(x, i): return pointer to node containing the i th smallest key of the subtree rooted at x .

OS-SELECT(x, i)

- $r \leftarrow \text{size}[\text{left}[x]] + 1$
- if $i = r$
- then return x
- elseif $i < r$
- then return OS-SELECT(left[x], i)
- else return OS-SELECT(right[x], $i - r$)
- Initial call: OS-SELECT(root[T], i)
- Example: OS-SELECT(root[T], 5) (see figure).

- Note: It is OK for keys to not be distinct. Rank is defined with respect to position in inorder walk. So if we changed D to C, rank of original C is 2, rank of D changed to C is 3.



Proof of Correctness and Efficiency

Correctness: r = rank of x within subtree rooted at x .

- If $i = r$, then we want x .
- If $i < r$, then i th smallest element is in x 's left subtree, and we want the i th smallest element in the subtree.
- If $i > r$, then i th smallest element is in x 's right subtree, but subtract off the r elements in x 's subtree that precede those in x 's right subtree.
- Like the randomized SELECT algorithm!

Analysis: Each recursive call goes down one level. Since R-B tree has $O(\lg n)$ levels, have $O(\lg n)$ calls $\Rightarrow O(\lg n)$ time.

Randomized Select

Order Statistics Chapter,
Finding the element of order i
REMINDER

- RANDOMIZED-SELECT(A, p, r, i)
- 1 **if** $p = r$
- 2 **then return** $A[p]$
- 3 $q \leftarrow \text{RANDOMIZED-PARTITION}(A, p, r)$
- 4 $k \leftarrow q - p + 1$
- 5 **if** $i = k$ \triangleright *the pivot value is the answer*
- 6 **then return** $A[q]$
- 7 **elseif** $i < k$
- 8 **then return** RANDOMIZED-SELECT($A, p, q - 1, i$)
- 9 **else return** RANDOMIZED-SELECT($A, q + 1, r, i - k$)

expected time of RANDOMIZED-SELECT is $\Theta(n)$.

OS-RANK(T, x)

//OS-RANK(T, x): return the rank of x in the linear order determined by an inorder walk of T .

- OS-RANK(T, x)
- $r \leftarrow \text{size}[\text{left}[x]] + 1$
- $y \leftarrow x$
- while $y \neq \text{root}[T]$
- do if $y = \text{right}[p[y]]$
- then $r \leftarrow r + \text{size}[\text{left}[p[y]]] + 1$
- $y \leftarrow p[y]$
- return r


Correctness

Loop invariant: At start of each iteration of while loop, r = rank of $\text{key}[x]$ in subtree rooted at y .

Initialization: Initially, r = rank of $\text{key}[x]$ in subtree rooted at x , and $y = x$.

Termination: Loop terminates when $y = \text{root}[T] \Rightarrow$ subtree rooted at y is entire tree. Therefore, r = rank of $\text{key}[x]$ in entire tree.

Maintenance: At end of each iteration, set $y \leftarrow p[y]$. So, show that if r = rank of $\text{key}[x]$ in subtree rooted at y at start of loop body, then r = rank of $\text{key}[x]$ in subtree rooted at $p[y]$ at end of loop body.



$[r = \# \text{ of nodes in subtree rooted at } y \text{ preceding } x \text{ in inorder walk}]$

Correctness (cont)


Maintenance: At end of each iteration, set $y \leftarrow p[y]$. So, show that if $r = \text{rank of key}[x]$ in subtree rooted at y at start of loop body, then $r = \text{rank of key}[x]$ in subtree rooted at $p[y]$ at end of loop body.

Must add nodes in y 's sibling's subtree.

If y is a left child, its sibling's subtree follows all nodes in y 's subtree \Rightarrow don't change r .

If y is a right child, all nodes in y 's sibling's subtree precede all nodes in y 's subtree \Rightarrow add size of y 's sibling's subtree, plus 1 for $p[y]$, into r .

Analysis: y goes up one level in each iteration \Rightarrow $O(\lg n)$ time.



Efficiency

Maintaining subtree sizes

Need to maintain *size[x]* fields during *insert and delete operations*.

Need to maintain them efficiently.


Otherwise, might have to recompute them all, at a cost of (n) .

Will see how to maintain without increasing $O(\lg n)$ time for *insert and delete*.

Efficiency

- ***Insert:***
- During pass downward, we know that the new node will be a descendant of
- each node we visit, and only of these nodes. Therefore, increment *size field* of each node visited.
- Then there is the fixup pass:
 - • Goes up the tree.
 - • Changes colors $O(\lg n)$ times.
 - • Performs ≤ 2 rotations.
 - • Color changes don't affect subtree sizes.
 - • Rotations do!
 - • But we can determine new sizes based on old sizes and sizes of children.

Efficiency (cont)



$size[y] \leftarrow size[x]$

$size[x] \leftarrow size[left[x]] + size[right[x]] + 1$

- Similar for right rotation.
- Therefore, can update in $O(1)$ time per rotation $\Rightarrow O(1)$ time spent updating
- size fields during fixup.
- Therefore, $O(\lg n)$ to insert.

Efficiency (cont)

- **Delete: Also 2 phases:**
 1. Splice out some node y .
 2. Fixup.
- After splicing out y , *traverse a path $y \rightarrow \text{root}$, decrementing size in each node on path. $O(\lg n)$ time.*
- During fixup, like insertion, only color changes and rotations.
- • ≤ 3 rotations $\Rightarrow O(1)$ time spent updating size fields during fixup.
- • Therefore, $O(\lg n)$ to delete.
- Done!

Methodology for augmenting a data structure


- 1. Choose an underlying data structure.
 - 2. Determine additional information to maintain.
 - 3. Verify that we can maintain additional information for existing data structure operations.
 - 4. Develop new operations.
- 1. R-B tree.
 - 2. *size[x]*.
 - 3. Showed how to maintain *size during insert and delete*.
 - 4. Developed OS-SELECT and OS-RANK.

Red-Black Trees Amenable to Augmentation

- **Theorem**
- Augment a R-B tree with field f , where $f[x]$ depends only on information in x , $\text{left}[x]$, and $\text{right}[x]$ (including $f[\text{left}[x]]$ and $f[\text{right}[x]]$). Then can maintain values of f in all nodes during insert and delete without affecting $O(\lg n)$ performance.
- **Proof** Since $f[x]$ depends only on x and its children, when we alter information in x , changes propagate only upward (to $p[x]$, $p[p[x]]$, . . . , root).
- Height = $O(\lg n) \Rightarrow O(\lg n)$ updates, at $O(1)$ each.
- **Insertion: see the book**
- **Delete: see the book**

Interval Trees

Maintain a set of intervals. For instance, time intervals.



Interval Tree Properties

- **Operations**
 - $\text{INTERVAL-INSERT}(T, x)$: $\text{int}[x]$ already filled in.
 - $\text{INTERVAL-DELETE}(T, x)$
 - $\text{INTERVAL-SEARCH}(T, i)$: return pointer to a node x in T such that $\text{int}[x]$ overlaps interval i . Any overlapping node in T is OK. Return pointer to sentinel $\text{nil}[T]$ if no overlapping node in T .
- Interval i has $\text{low}[i]$, $\text{high}[i]$.
- i and j overlap if and only if $\text{low}[i] \leq \text{high}[j]$ and $\text{low}[j] \leq \text{high}[i]$.

Augmenting to Get Interval-Trees

For interval trees

1. Use R-B trees.


Each node x contains interval $int[x]$.

Key is low endpoint ($low[int[x]]$).

Inorder walk would list intervals sorted by low endpoint.

2. Each node x contains

$max[x] = \text{max endpoint value in subtree rooted at } x$.



Summary

Binary Search Tree (BST) review

Red and Black Trees

2-3 and 2-3-4 trees

Operations on Red and Black Trees